

CIS 5800

# Machine Perception

Instructor: Lingjie Liu  
Lec 25: April 30, 2025

# Administrivia

Homework 5 (optional, 6 points for grade compensation) has been released. It and the small projects are due on May 14.

Final exam coming up

- Date reminder: Wednesday May 7, 3-5pm in DRLB A1. (Info is on [courses.upenn.edu](https://courses.upenn.edu))
- Syllabus: Mainly the material covered in class after Wed March 19 (not covered by mid-term exam).
- Review lecture in the last class on Wed April 30.
- If you are unable to attend the midterm exam in person on May 7, please complete the form by April 30: <https://forms.gle/JwaAxrKGfBzoo6z77>
- Also, you need to contact the [Weingarten Office](#) for academic accommodations and send me the paperwork or approval from the Weingarten Office.

Putting the class in perspective in the context of computer vision  
and a quick overview of “visual recognition”


# What Info can be Extracted from Images?



Source: S. Lazebnik

# What Info can be Extracted from Images?


This class!  
geometric  
information



Source: S. Lazebnik


# What Info can be Extracted from Images?

This class!  
geometric  
information




“visual  
recognition”  
semantic  
information

# ML in Computer Vision



**The *very old*: 1960's - Mid 1990's**

Image → hand-def. features → hand-def. classifier






**The *old*: Mid 1990's – 2012**

Image → hand-def. features → learned classifier


# What Should Good Visual Features Do?

What is a “good” feature space?



Good features make useful tasks easy to perform.

# What Should Good Visual Features Do?



How should we produce such good features?

# Most Feature Extraction Frameworks Pre-2012

**Step 1: Focus on “interest points” rather than all pixels**

E.g. corner points, “difference of gaussians”, or even a uniform grid

**Step 2: Compute features at interest points.**

E.g. “SIFT”, “HOG”, “SURF”, “GIST”, etc.

**Step 3: Convert to fixed-dimensional feature vector by measuring statistics of the features such as histograms**


E.g. “Bag of Words”, “Spatial Pyramids”, etc.




Bag-of-Words histogram

See libraries like VLFeat and OpenCV

Use your favorite ML model now!




# Machine Learning for Semantic Computer Vision




**The very old: 1960's - Mid 1990's**

Image → hand-def. features → hand-def. classifier



**The old: Mid 1990's – 2012**

Image → hand-def. features → learned classifier



**The new: 2012 – ?**


Image → jointly learned features + classifier with  
“deep” multi-layer neural networks

# “Deep” Learning


“Deep” multi-layer neural networks are **representation learners**.

Every layer improves upon its preceding layer, tailoring the representation to the task.

Image




$D$ -length  
feature  $x$



“dog”


# Some sample applications of semantic vision

[Girshick et al. CVPR14]




Object detection

[Toshev et al. CVPR14]



Pose detection (regression)


[Long et al. CVPR15]



Semantic segmentation

# Some sample applications of semantic vision

[Chopra et al. CVPR05]



Similarity metric learning

[Dosovitskiy et al. CVPR15]





Image generation

[Dong et al. ECCV 2014]



Low-level image processing:  
(superresolution, deblurring,  
image quality etc.)


# Game playing from visual inputs

CNN + Reinforcement learning

[Mnih et al, Nature '15]



[Silver et al, Nature '16]



# Generating art



See if you can tell  
artists' originals  
from machine  
style imitations at:  
<http://turing.deeprt.io/>

Paper: [Gatys et al, "Neural ... Style", arXiv '15](#)  
Code (torch): <https://github.com/jcjohnson/neural-style>

# Where to learn more about ML and semantic vision?

## Machine learning courses:

CIS 519, 520, 522 usually cover semantic computer vision briefly, as an application domain for machine learning techniques

## CIS 581 Computer Vision & Computational Photography

The basics of image processing and semantic computer vision.

## CIS 680 Advanced Machine Perception

Cutting-edge techniques in semantic (largely) computer vision, best taken after some introduction to ML.


## CIS 7000 Advanced Topics



# There is lots of ML in geometric vision too!

(The following slides are based on materials from Ben Mildenhall, Vincent Sitzmann and Stephen Lombardi)

# Neural Scene Representation and Neural Rendering




# **Neural Radiance Fields (NeRF)**

Mildenhall, Srinivasan, Tancik, Barron, Ramamoorthi, Ng  
ECCV 2020

# Neural Volumetric Rendering

# Neural Volumetric **Rendering**

querying the radiance value  
along rays through 3D space



*What colour?*


# Neural Volumetric Rendering

continuous, differentiable  
rendering model without  
concrete ray/surface intersections



# Neural Volumetric Rendering

using a neural network as a  
scene representation, rather  
than a voxel grid of data





Inputs: sparse, unstructured  
photographs of a scene




Outputs: representation allowing us to  
render *new* views of that scene

# Overview

- ▶ Volumetric rendering math
- ▶ Neural networks as representations for spatial data
- ▶ Neural Radiance Fields (NeRF)

# Traditional volumetric rendering



- ▶ Theory of volume rendering co-opted from physics in the 1980s: absorption, emission, out-scattering/in-scattering
- ▶ Adapted for visualising medical data and linked with alpha compositing
- ▶ Modern path tracers use sophisticated Monte Carlo methods to render volumetric effects

Chandrasekhar 1950, *Radiative Transfer*

Kajia 1984, *Ray Tracing Volume Densities*


Levoy 1988, *Display of Surfaces from Volume Data*

Max 1995, *Optical Models for Direct Volume Rendering*

Porter and Duff 1984, *Compositing Digital Images*

Novak et al 2018, *Monte Carlo methods for physically based volume rendering*

# Traditional volumetric rendering



Medical data visualisation  
[Levoy]

- ▶ Theory of volume rendering co-opted from physics in the 1980s: absorption, emission, out-scattering/in-scattering
- ▶ Adapted for visualising medical data and linked with alpha compositing
- ▶ Modern path tracers use sophisticated Monte Carlo methods to render volumetric effects

Alpha compositing [Porter and Duff]

Chandrasekhar 1950, *Radiative Transfer*

Kajia 1984, *Ray Tracing Volume Densities*

Levoy 1988, *Display of Surfaces from Volume Data*

Max 1995, *Optical Models for Direct Volume Rendering*

Porter and Duff 1984, *Compositing Digital Images*

Novak et al 2018, *Monte Carlo methods for physically based volume rendering*

# Traditional volumetric rendering



Physically-based Monte Carlo rendering [Novak et al]

- ▶ Theory of volume rendering co-opted from physics in the 1980s: absorption, emission, out-scattering/in-scattering
- ▶ Adapted for visualising medical data and linked with alpha compositing
- ▶ Modern path tracers use sophisticated Monte Carlo methods to render volumetric effects

Chandrasekhar 1950, *Radiative Transfer*

Kajia 1984, *Ray Tracing Volume Densities*


Levoy 1988, *Display of Surfaces from Volume Data*

Max 1995, *Optical Models for Direct Volume Rendering*

Porter and Duff 1984, *Compositing Digital Images*

Novak et al 2018, *Monte Carlo methods for physically based volume rendering*

# Volumetric rendering and machine learning



- ▶ Various volume-rendering-esque methods devised for 3D shape reconstruction methods
- ▶ Scaled up to higher resolution volumes to achieve excellent view synthesis results

Tulsiani et al 2017, *Multi-view Supervision for Single-view Reconstruction via Differentiable Ray Consistency*

Henzler et al 2019, *Escaping Plato's Cave: 3D Shape From Adversarial Rendering*

Zhou et al 2018, *Stereo Magnification: Learning View Synthesis using Multiplane Images*

Lombardi et al 2019, *Neural Volumes: Learning Dynamic Renderable Volumes from Images*

# Volumetric rendering and machine learning



Slices from a volumetric scene representation [Zhou et al]

- ▶ Various volume-rendering-esque methods devised for 3D shape reconstruction methods
- ▶ Scaled up to higher resolution voxel grids, ML methods can achieve excellent view synthesis results



View synthesis from a dynamic voxel grid [Lombardi et al]


Tulsiani et al 2017, *Multi-view Supervision for Single-view Reconstruction via Differentiable Ray Consistency*

Zhou et al 2018, *Scanning Holes have Everything to do with Volumetric Rendering Images*

Lombardi et al 2019, *Neural Volumes: Learning Dynamic Renderable Volumes from Images*


# Volumetric formulation for NeRF

# Volumetric formulation for NeRF




Scene is a cloud of tiny colored particles

# Volumetric formulation for NeRF




# Volumetric formulation for NeRF



This notion is *probabilistic*: chance that ray stops in a small interval around  $t$  is  $\sigma(t) dt$ .

$\sigma(t)$  is known as the “volume density”


# Volumetric formulation for NeRF



To determine if  $t$  is the *first* hit, need to know  $T(t)$ :  
probability that the ray didn't hit any particles earlier.

$T(t)$  is called “transmittance”

# Volumetric formulation for NeRF




To determine if  $t$  is the *first* hit, need to know  $T(t)$ :  
probability that the ray didn't hit any particles earlier.

$T(t)$  is called “transmittance”


We assume  $\sigma$  is known and want to use it to calculate  $T(t)$

# Volumetric formulation for NeRF



$\sigma$  and  $T$  are related by the probability fact that  
 $P[\text{no hits before } t + dt] = P[\text{no hits before } t] \times P[\text{no hit at } t]$

# Volumetric formulation for NeRF



These are related by the probability fact that

$$T(t + dt) = T(t) \times (1 - \sigma(t)dt)$$

# Volumetric formulation for NeRF

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

# Volumetric formulation for NeRF

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Split up differential  $\Rightarrow$   $T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$

# Volumetric formulation for NeRF

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Split up differential  $\Rightarrow$   $T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$

Rearrange  $\Rightarrow$   $\frac{T'(t)}{T(t)}dt = -\sigma(t)dt$

# Volumetric formulation for NeRF

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

Split up differential  $\Rightarrow$   $T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt$

Rearrange  $\Rightarrow$   $\frac{T'(t)}{T(t)}dt = -\sigma(t)dt$

**Integrate**  $\Rightarrow$   $\log T(t) = - \int_{t_0}^t \sigma(s)ds$

$$\Rightarrow T(t) = \exp\left(- \int_{t_0}^t \sigma(s)ds\right)$$

# Volumetric formulation for NeRF

Thus, the probability that a ray first hits a particle at  $t$  is

$$T(t)\sigma(t) dt = \exp\left(-\int_{t_0}^t \sigma(t) dt\right) \sigma(t) dt$$

# Volumetric formulation for NeRF

Thus, the probability that a ray first hits a particle at  $t$  is


$$T(t)\sigma(t) dt = \exp\left(-\int_{t_0}^t \sigma(s) ds\right) \sigma(t) dt$$

And expected color returned by the ray will be

$$\int_{t_0}^{t_1} T(t)\sigma(t)\mathbf{c}(t) dt$$


Note the nested integral!

# Approximating the nested integral




We use quadrature to approximate the nested integral,

# Approximating the nested integral




We use quadrature to approximate the nested integral,  
splitting the ray up into  $n$  segments with endpoints  $\{t_1, t_2, \dots, t_{n+1}\}$

# Approximating the nested integral



We use quadrature to approximate the nested integral,  
splitting the ray up into  $n$  segments with endpoints  $\{t_1, t_2, \dots, t_{n+1}\}$   
with lengths  $\bar{\delta}_i = t_{i+1} - t_i$

# Approximating the nested integral



We assume volume density and color are  
roughly constant within each interval

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt$$

This allows us to break the outer integral

# Approximating the nested integral

$$\int T(t) \sigma(t) \mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t) \sigma_i \mathbf{c}_i dt$$

This allows us to break the outer integral into a sum of analytically tractable integrals

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

Catch: piecewise constant density and color  
**do not** imply constant transmittance!

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t) \boxed{\sigma(t)} \mathbf{c}_i dt$$

Catch: piecewise constant density and color  
**do not** imply constant transmittance!

Important to account for how early part of a segment blocks later part when  $\sigma_i$  is high

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

For  $t \in [t_i, t_{i+1}]$ ,  $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

For  $t \in [t_i, t_{i+1}]$ ,  $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$

  $\exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right) = T_i$  "How much is blocked by all previous segments?"

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

For  $t \in [t_i, t_{i+1}]$ ,  $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$

"How much is blocked partway  
through the current segment?"



$$\exp(-\sigma_i(t - t_i))$$

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

# Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

Substitute

$$= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt$$

# Approximating the nested integral

$$\begin{aligned} \int T(t)\sigma(t)\mathbf{c}(t) dt &\approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt \\ &= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt \\ \text{Integrate} \quad &= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i} \end{aligned}$$

# Approximating the nested integral

$$\begin{aligned} \int T(t)\sigma(t)\mathbf{c}(t) dt &\approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt \\ &= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt \\ &= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i} \\ \text{Cancel } \sigma_i &\quad = \sum_{i=1}^n T_i\mathbf{c}_i (1 - \exp(-\sigma_i\delta_i)) \end{aligned}$$

# Summary: volume rendering integral estimate

Rendering model for ray  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ :

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$


↑  
weights      colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment  $i$ :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$



# Summary: volume rendering integral estimate

Rendering model for ray  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ :

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$


↑  
weights      colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment  $i$ :


$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$



# Overview


- ▶ Volumetric rendering math
- ▶ Neural networks as representations for spatial data
- ▶ Neural Radiance Fields (NeRF)

# Toy problem: storing 2D image data



Usually we store an image as a  
2D grid of RGB color values

# Toy problem: storing 2D image data




What if we train a simple fully-connected  
network (MLP) to do this instead?

# Naive approach fails!

Ground truth image



Standard fully-connected net



## Problem:

“Standard” coordinate-based MLPs cannot represent high-frequency functions

Solution:

Pass input coordinates through a  
high frequency mapping first

# Input coordinate mapping

- ▶ Simple formula: apply a tall skinny matrix  $\mathbf{B}$  to input coordinate vector  $\mathbf{x}$ , then pass through  $\sin$  and  $\cos$ :

$$\gamma(\mathbf{x}) = (\sin(2\pi \mathbf{B}\mathbf{x}), \cos(2\pi \mathbf{B}\mathbf{x}))$$

- ▶ Passing network a subset of the Fourier basis functions. Same effect from:
  - ▶ Positional encoding
  - ▶ Fourier features
  - ▶ SIREN

# Problem solved


Ground truth image



Standard fully-connected net




With “positional encoding”




# Overview

- ▶ Volumetric rendering math
- ▶ Neural networks as representations for spatial data
- ▶ Neural Radiance Fields (NeRF)

*NeRF* = volume rendering +  
coordinate-based network



# Train network to reproduce input views of scene using gradient descent



# Visualizing view-dependent effects



Regular NeRF rendering



Manipulating input viewing directions

# Visualizing learned density field as geometry



Regular NeRF rendering



Expected ray termination depth

# Visualizing learned density field as geometry



Regular NeRF rendering



Expected ray termination depth

If you're interested, you may take: CIS 7000-005 Introduction to Neural Scene Representation and Neural Rendering, in Fall 2025.